

AI ECONOMICS / ENTERPRISE DATA

Why the numbers on the AI revolution don't add up.

The sharpest scepticism about AI economics is now coming from inside the world's largest institutions. J.P. Morgan analysts have modelled the AI infrastructure buildout and concluded that earning a 10% return on it would require roughly \$650 billion of annual revenue — in perpetuity. Hyperscaler capex passed \$400 billion in 2025 and is heading toward \$700 billion. Even JPMorgan's own celebrated AI programme spends roughly \$2 billion a year to generate a comparable amount of value — break-even, at a scale most banks cannot match. The institutions writing the largest cheques are asking the hardest question: where does the return come from?

WHY THE MATHS BREAKS

01

Compute demand compounds, revenue doesn't

Reasoning models can need over 100x the compute of single-shot inference, and one agentic workflow can trigger dozens of calls. Falling unit prices drive heavier usage — total spend keeps rising even as each query gets cheaper.

02

The spend is variable, the value is not

Compute bills scale with every pipeline, reload and experiment. Business value arrives in a handful of proven use cases. When cost grows linearly and value doesn't, the gap widens with adoption — the opposite of what the board was promised.

03

Most compute feeds preparation, not intelligence

Years of platform and pipeline work precede any AI value — JPMorgan built entire data and governance platforms before deployment. In most enterprises the majority of "AI spend" is really data ingestion, cleaning and reprocessing, billed by the second.

04

Bad data multiplies the bill

Every ungoverned dataset consumed by a model is compute spent twice: once ingesting it, again correcting what it produced. AI built on untrusted data doesn't just fail to return — it actively compounds the cost of failure.

Making the numbers add up: fix the denominator

You cannot control what hyperscalers spend — but you can control your own equation. The return on AI is value divided by the cost of trusted data, and that denominator is where QBricks works. Automated, streaming ingestion moves only incremental change, ending the full reloads that dominate compute bills. Every record is validated against an ODCS data contract before it lands — registered in the governance catalog, auditable before and after, with a human in the loop — so models consume trusted data instead of burning compute on garbage in, garbage out. And because it runs natively on Databricks, Microsoft Fabric and Snowflake with local compute options, heavy ingestion work can run at fixed cost on infrastructure you own.

The problem is not the AI solution you are trying to implement. It is the quality of your data.

QBricks: automated, governed, incremental ingestion — from system of record to lakehouse in hours.

[Request a demo →](#)