

FINANCE / DATA PLATFORM ECONOMICS

The CFO is waking up to the cost of compute.

Cloud data platforms were sold on elasticity: pay only for what you use. The CFO is now discovering what "what you use" actually means — a variable cost that grows with every pipeline, every reload and every new AI initiative, with no natural ceiling. Data compute is becoming one of the fastest-growing operating expenses in the business, and unlike headcount or licences, nobody signed off on it. When the unit cost of serving the business rises faster than revenue, the business model itself is at risk.

WHERE THE VARIABLE COST COMES FROM

The driver is rarely the analytics the business consumes. It is the machinery behind it: hand-built pipelines that reload entire tables nightly to capture a day's worth of change; bad data reprocessed three layers deep before anyone notices; thousands of ungoverned notebooks, each quietly billing by the second; and duplicate copies of the same source ingested by teams who couldn't find — or couldn't trust — the first one. Every one of these burns compute to produce nothing new. The spend is not the price of insight; it is the price of ungoverned ingestion.

FOUR STRATEGIES TO CAP IT — AND HOW QBRICKS DELIVERS EACH

01

Pay for change, not for volume

QBricks streams incremental change from the system of record. Compute scales with what actually changed today — not with the size of tables you already hold. The nightly full reload, the largest single line on most compute bills, disappears.

02

Reject bad data before it costs anything

Every record is validated against an ODCS data contract at ingestion, with a human in the loop for exceptions. Defects are stopped at the door instead of being reprocessed through Bronze, Silver and Gold at your expense.

03

Replace notebooks with one governed factory

Automated ingestion replaces thousands of hand-built pipelines — and the engineering armies that maintain them. One auditable process, one place to measure cost, ingested once and reused by everyone.

04

Put compute where it's cheapest

QBricks runs natively on Databricks, Microsoft Fabric and Snowflake — with local compute options for heavy ingestion work. Validation can run on infrastructure you own at fixed cost, so premium platform compute is reserved for the analytics that earn revenue.

A different shape of cost curve

The financial outcome is a change in the shape of the cost curve, not just its level: ingestion becomes a governed, predictable process whose cost tracks business change — auditable line by line, before and after — instead of an open-ended variable spend that compounds with every new use case.

Make compute a controlled cost, not a variable one.

QBricks: automated, governed, incremental ingestion — from system of record to lakehouse in hours.

[Request a demo →](#)