

Compute cost outlook: the commodity turn

Compute is becoming a commodity. Hedge it like one.

On 5 October 2026 the first US-regulated compute futures list on NYMEX. Once a forward curve exists, every data platform's compute bill becomes a marked exposure. This paper explains how the market is forming, why a paper hedge alone leaves the lakehouse uncovered, and how QBricks provides the operational hedge: consume a fraction of the commodity in the first place.

5 Oct 2026

CME lists H100 and B200 rental-index futures on NYMEX — compute's first exchange-cleared forward curve

+56%

move in H100 one-year contract rates between October 2025 and March 2026 — the tenor enterprises actually buy

€0.48

QBricks compute per TB ingested, measured 29 September 2026 — a cost the curve cannot reprice

Executive summary

Compute has spent two decades as an operating expense buried in an infrastructure budget. In 2026 it is crossing into the category of oil, power and freight: a priced input with an index, a forward curve, cleared futures and a regulator consulting on how to supervise them. CME Group, Intercontinental Exchange and Architect Financial Technologies have each announced contracts; the CFTC opened its consultation in August; the first listing is days away.

For data platform leaders this changes the question. Today's lakehouse compute — DBUs, DPUs, executor-hours — is bought in a fragmented, negotiated market with no visible reference price. Once a benchmark exists, suppliers index to it, finance marks against it, and a bill that used to be "what the cloud charged us" becomes an exposure a treasury expects to see managed.

Futures help, but only partly. They reference GPU-hours, not the vCPU-plus-DRAM units Spark bills in; they fix a price but not a quantity; they carry basis risk, margin and roll cost; and if the curve falls, the hedge loses money. The durable hedge for any commodity is the one airlines learned with jet fuel: swaps for the short end, a more efficient fleet for the long run.

QBricks is the efficient fleet. It keeps the lakehouse as the governed landing zone and routes roughly 70 percent of workloads onto EOS, a Spark-free vectorised engine on one right-sized VM, at a measured compute cost of €0.48 per terabyte. Modelled against a volatile compute index, the routed estate runs 47–59 percent below the unhedged cost every year — including the years in which a futures hedge alone goes underwater. Paid by gain-share, the hedge funds itself.

From procurement line item to exchange-cleared contract

Bilateral compute forwards have existed for years: every neocloud term agreement and every reserved-instance commitment is a forward, negotiated over the counter with no central clearing. What is new in 2026 is the migration of that exposure into standardised, cash-settled, exchange-cleared contracts denominated in dollars per GPU-hour — the step that admits institutional capital and produces public price discovery.

Three exchange-level announcements landed within four months. Each follows the energy-market template: an independent index built from printed transactions, a monthly contract, financial settlement, no physical delivery.

Timeline: the financialisation of compute, 2026

Jan	Architect Financial Technologies announces perpetual futures on GPU rental and DRAM, benchmarked to Ornn index data.
12 May	CME Group and Silicon Data announce a GPU compute futures market, pending CFTC review.
19 May	Intercontinental Exchange and Ornn announce a competing suite of USD-denominated, cash-settled futures on the Ornn Compute Price Index.
11 Aug	CME confirms two NYMEX contracts — Silicon Data H100 and B200 Rental Index Futures — each one month of hourly GPU rental, listing 5 October.
19 Aug	The CFTC opens a public consultation on AI compute derivatives: financialisation moves from corporate experiment to regulatory agenda.
5 Oct	First US-regulated compute futures trade. Compute has a public forward curve.

Anatomy of one compute future

1 contract = 1 month of 1 GPU

Quoted in dollars and cents per GPU-hour; a one-cent tick is worth \$7.30, so each contract covers roughly 730 GPU-hours. Cash-settled against the index — no hardware ever changes hands. Separate contracts per chip: an H100-hour and a B200-hour are different products, not a blended unit of compute.


Why a megawatt-hour is easy and a GPU-hour is not

Commodity markets run on rude simplifications: a barrel of WTI stands in for oil, a TEU for freight. Compute is semi-fungible — interconnect, storage pipeline and scheduler change how much work an hour delivers — and it is non-storable: an idle hour is gone, so prices move on utilisation and opportunity cost, not inventory. The standard commodity toolkit does not transfer cleanly. The market is forming anyway.

What a forward curve does to a lakehouse budget

A benchmark changes behaviour on both sides of a trade. Suppliers index to it: once there is a printed H100 rate, every cloud and neocloud price list drifts toward it, as European cloud pricing already moved with memory in 2026. Buyers are marked against it: the moment a reference price is public, a CFO can see what the platform’s compute would have cost at the curve, and a variance that used to be invisible becomes a line in the risk report.

The curve today is in mild backwardation — the market expects GPU rents to ease as Blackwell volume lands — but the path there is anything but smooth. The one-year H100 rate rose 56 percent in the five months to March 2026. A data platform whose Spark fleet reprices on a lag to that market is carrying commodity volatility without a commodity desk.



Price discovery

A public reference replaces fragmented, negotiated quotes. Variance from the curve becomes visible to finance.

Indexation

Suppliers reprice to the benchmark. Managed Spark units — vCPU plus a fixed parcel of DRAM — follow with a lag.

Mark-to-market

Compute joins fuel and FX as an exposure the treasury expects to see hedged, reported and audited.

Why futures alone leave the lakehouse uncovered

Compute futures are a genuine advance. A platform that knows it will buy GPU capacity next year can lock a price today, and a data-centre owner can lock revenue. But a lakehouse compute bill is a poor fit for the contracts as listed, for five reasons.

Basis. The contracts reference H100 and B200 rental rates. Spark bills in vCPU-plus-DRAM units. The two move together — both are priced off the same scarce memory and data-centre capacity — but not one for one. A hedge in the wrong underlying is a correlation bet.

Direction. The curve is backwardated. Lock a forward and the market falls, and the hedge loses — a real cash cost that a treasury must explain while the platform team’s budget looks fine.

Liquidity. Day-one depth is unknown. Standardisation, index reliability and the perishable nature of the underlying are open questions the regulator itself has raised.

Price, not quantity. A future fixes the price of a unit. It does nothing about how many units a Spark estate consumes — shuffle, serialisation, idle autoscale headroom and per-step minimums are all bought at the locked price.

Carry. Monthly contracts must be rolled; cleared positions post daily variation margin. Hedging three years of compute means a standing programme, not a single trade.

What it does solve. Visibility. A listed curve ends the era of compute as an unexamined expense — which is exactly why the operational hedge is about to be measured.

Two hedges, one exposure

	Financial: compute futures	Operational: QBricks
What it fixes	Price per GPU-hour, for the contract month	Quantity consumed — ~70% of workloads leave the metered cluster
Underlying	H100 / B200 rental index — basis risk to Spark units	Your own workloads — no basis
If the curve falls	Hedge loses money	Saving shrinks, fee shrinks with it — never negative
Tenor	Monthly, rolled; margin posted daily	Structural — persists for the life of the estate
Cost of the hedge	Spread, roll, margin funding	Gain-share on audited savings — self-funding

The two are complementary, not rivals: futures for the residual Spark share and the GPU estate; QBricks for the volume that need never touch the curve.

Data lands in the lakehouse. QBricks runs the workload.

An airline does not hedge jet fuel only with swaps. It buys a more efficient fleet, because efficiency is the one hedge that pays in every price scenario. The data-platform equivalent is to stop running most of the estate on an engine shaped for cheap memory and expensive coordination.

Spark parallelises across many executors, each holding a slice of the dataset in RAM, and bills per executor-hour. EOS, the SQL engine behind QBricks, inverts that trade: a hardened fork of Apache DataFusion with Arrow-native streaming and the Vortex columnar format, built in Rust, running an organisation’s pipeline estate on one right-sized VM. No cluster, no shuffle, no idle headroom, no memory tax — and a compute line that is a single committed item, forecastable to the euro.

QBricks puts EOS behind its governed ingestion layer. Databricks, Microsoft Fabric or Snowflake remain the governed landing zone on Iceberg or Delta tables. Each workload is routed: small scans, simple aggregates and incremental merges to EOS on a single node; large joins, streaming and ML to Spark where it is the right tool. Every record is still checked against an Open Data Contract Standard contract, with a human in the loop and before-and-after files for audit.

Measured on 29 September 2026: Azure, production code path

12.08 _{TB}	50.6 _{s/TB}	€0.48 _{/TB}	8.1 _{vCPU·h}
CSV ingested · 86.6 billion rows · 8/8 tables validated	Sustained 18.42 GiB/s · 71 TB an hour	Timed-ingestion compute · €3.74 incl. all observed compute and network	Worker compute per TB — the quantity the curve would price

Single successful run on a TPC-H-derived SF10000 dataset; nine Standard_D64ds_v5 workers. Cost per TB is what makes the hedge legible to finance: a unit price, measured and auditable, that can be set beside the curve.

Paper hedge, operational hedge, both

Take a bank running twenty source systems into a Databricks lakehouse, Spark compute indexed to 100 in September 2025. Assume a volatile compute index that spikes to 160 in 2027 before easing with new supply. Three strategies: lock 70 percent of consumption at the forward curve (115) with 10 percent basis slippage; route 70 percent of workloads to QBricks at 25 percent of the Spark cost in 2027–28 and 30 percent thereafter; or do both.

Year	Unhedged all Spark	Futures only 70% locked at 115	QBricks only 70% routed	Both
2027	160	137 -15%	76 -53%	66 -59%
2028	150	134 -11%	71 -53%	64 -57%
2029	130	128 -2%	66 -49%	65 -50%
2030	120	125 +4%	61 -49%	63 -47%

Illustrative; index Sep 2025 = 100. Futures: $0.7 \times 115 \times 1.10 \text{ basis} + 0.3 \times \text{spot}$. QBricks: $0.7 \times \text{spot} \times 25\% \text{ (30\% from 2029)} + 0.3 \times \text{spot}$. Both: QBricks on the routed share, futures on the residual Spark share. Real pass-through is lagged and partly cushioned by long-term agreements; the relative gap is the point. The 70% routed share is an example — every additional 10 points routed moves the blended index roughly 8% closer to the QBricks line.

Reading the table

The futures hedge does its job in the spike year and then decays as the curve falls below the locked price, turning negative by 2030. The operational hedge is indifferent to direction: it removes volume, so it saves roughly half the bill whether the index is 160 or 120. Combining the two adds a few points in the spike years and nothing after — the residual Spark share is small enough that the paper hedge barely matters.

Commercial model: gain-share

QBricks is paid from realised savings, not a licence that grows with the cluster. Compute is measured before and after routing on the same workloads, the comparison files are retained for audit, and the agreed share is invoiced against the verified difference. If the curve flattens sooner than modelled, saving and fee shrink together. The hedge is self-funding in every scenario — the one property no futures position can offer.

Own the volume, not just the price

Compute is becoming a commodity because the AI build-out made it scarce enough to be worth pricing. The exchanges, the index providers and the regulator are building the market whether data platforms are ready or not — and once the curve prints, every unhedged compute budget will be visible against it.

A treasury's answer to a commodity exposure has two parts: hedge the price where instruments exist, and reduce the quantity wherever engineering allows. Futures cover the first for the GPU estate. QBricks covers the second for the lakehouse — keep the landing zone, keep the governance, drop the cluster for most of the work, and pay for the hedge out of what it saves.

Request a routing assessment.

Two weeks, on your live workload mix, no pipeline build. It returns the routed share, the blended index against the curve and the audit files the gain-share contract runs on.

qbricks.ai

Sources and notes

1. CME Group press releases, 12 May and 11 Aug 2026: Silicon Data H100 and B200 Rental Index Futures, NYMEX listing 5 Oct 2026 pending regulatory review; contract specification (one month of hourly rental, \$0.01 tick = \$7.30, cash-settled).
2. Intercontinental Exchange, 19 May 2026: USD-denominated, cash-settled futures on the Ornn Compute Price Index. Architect Financial Technologies, Jan 2026: GPU and DRAM perpetual futures on Ornn benchmark data.
3. CFTC request for public comment on AI compute derivatives, 19 Aug 2026 (Bloomberg).
4. Silicon Data, "GPU Futures" (mid-2026): spot and 36-month term rates for B200, H100, A100; H100 one-year contract rate from ~\$1.70/hr (Oct 2025) to ~\$2.65/hr (Mar 2026).
5. Assody, "Pricing Compute Futures: Forward Curve and Volatility for a Non-Storable, Depreciating Commodity", SSRN 2026, on the limits of the standard commodity toolkit.
6. Memory pass-through and lakehouse billing units: QBricks, "QBricks: a hedge to the cost of rising memory, Spark and compute", Sep 2026 — TrendForce contract-price data; OVHcloud pricing notice; AWS Glue DPU definition.
7. QBricks ingest benchmark: 12.082 TB / 86.6 bn rows in 610.80 s on 9 × Azure Standard_D64ds_v5, €5.79 timed compute (€0.48/TB), €45.20 combined observed compute and network (€3.74/TB); 29 Sep 2026, single run. Full report available on request.
8. Worked example and routing shares: Infinium Consulting illustrative model, Oct 2026. Forecasts, not actuals; depends on the client's workload mix. Not investment advice.