

PLATFORM ARCHITECTURE / COMPUTE ECONOMICS

A New Paradigm In Data Management.

The prevailing answer to slow data platforms has been the same for a decade: add nodes. Bigger Spark clusters, larger capacities, more parallelism — infrastructure scaled to survive the nightly peak of reloading everything, then sitting idle the rest of the day. QBricks starts from the opposite premise: the work itself is wrong. Move only what changed, validate it before it lands, and the cluster you were about to double becomes a cluster you no longer need.

Where the servers go

01 Streaming replaces the batch peak

Clusters are sized for the worst hour of the night — the full reload window. QBricks streams incremental change continuously, so there is no peak to size for. A small, steady flow needs a fraction of the nodes a nightly avalanche does.

02 Validation runs on local compute

Record-level contract validation — the heavy, repetitive work of ingestion — doesn't need premium lakehouse clusters. QBricks runs it on local compute you already own, at fixed cost, applying each platform's security standards. Cloud nodes are reserved for the analytics that earn revenue.

03 One process replaces thousands of pipelines

Every hand-built notebook carries its own cluster, schedule and idle time. One automated, governed ingestion process consolidates that sprawl — fewer jobs, fewer clusters, no duplicate copies of the same source ingested by teams who couldn't find the first.

04 Bad data never reaches paid compute

Records that fail their ODCS data contract are stopped at ingestion, with a human in the loop. Downstream nodes never process, store or re-correct data that was never fit for use — work that today silently inflates every cluster.

The paradigm shift

The old model scales infrastructure to tolerate ungoverned work: more nodes, bigger SKUs, taller bills. The new model governs the work so the infrastructure can shrink: incremental by default, validated at the point of entry, ingested once and reused everywhere, with the heaviest processing on the cheapest compute. QBricks runs this natively on Databricks, Microsoft Fabric and Snowflake — auditable end to end, before and after files, agents isolated from the data lake. The result is a platform whose footprint tracks business change, not data volume: fewer nodes, fewer servers, smaller capacities, and a compute bill that finally has a ceiling.

Don't scale the cluster. Shrink the work.

QBricks: automated, governed, incremental ingestion — from system of record to lakehouse in hours.

[QBRICKS.AI](https://qbricks.ai)