

BANKING / DATA ENGINEERING

Are You Stuck At The Bronze Layer When Implementing Databricks?

Most banks adopt the medallion architecture with confidence: land raw data in Bronze, refine it to Silver, serve trusted Gold. Then reality arrives. Two years in, the lakehouse holds terabytes of raw extracts, a handful of hand-built pipelines — and almost nothing the business can trust. The problem is not Databricks. It is how the data gets there.

Five reasons banks stall

01 Hand-built pipelines don't scale

A bank has hundreds of systems of record. Building each pipeline by hand — one notebook, one engineer, one quarter at a time — means the backlog grows faster than the team can deliver.

02 Governance is bolted on, not built in

BCBS 239, DORA and GDPR demand lineage, quality evidence and auditability. When ungoverned data lands first and controls come later, every dataset needs remediation before it can leave Bronze.

03 Source knowledge lives in silos

Core banking, payments and risk systems carry decades of undocumented semantics. Without agreed definitions, Silver becomes a negotiation between engineering and the business — and it rarely ends.

04 Batch thinking in a streaming world

Nightly full loads reprocess everything to change almost nothing. Costs climb, SLAs slip, and incremental change — the thing regulators and ML models actually need — is lost in the reload.

05 Quality is discovered downstream

Defects surface in a regulatory report or an AI model — weeks after ingestion, at the most expensive point to fix. Without record-level validation at the point of entry, Bronze becomes a landfill, and every AI initiative built on it inherits the mess.

The way out: contract-first ingestion

Banks that escape Bronze invert the model. Instead of landing raw data and cleaning it later, they agree the rules first — as machine-readable data contracts under the Open Data Contract Standard (ODCS) — and validate every record against them at ingestion. Governance stops being remediation and becomes the pipeline itself: schema, quality thresholds, ownership and lineage are enforced before data lands, registered in Unity Catalog, with a human always in the loop for exceptions. Automated, streaming ingestion replaces the hand-built backlog, so incremental change flows from system of record to lakehouse in hours, not quarters.

The problem is not the AI solution you are trying to implement. It is the quality of your data.

QBricks ingests from system of record to governed lakehouse — automated, auditable, ODCS-native on Databricks.

QBRICKS.AI